

A Machine Learning Based Approach for Predicting Adult Personal Information

Lezhong Huang

Johns Hopkins University
500 W. University Pkwy, 14F
Baltimore, MD 21210, USA
huanglezhong@gmail.com

Yang Dai

Johns Hopkins University
104 W. University Pkwy, 5F
Baltimore, MD 21210, USA
ydai9@jhu.edu

Abstract

This document contains how to implement a machine learning algorithm. This algorithm is applied on predicting some aspects of personal information given a large training set. The sample is a combination of many aspects of personal information, so any one of them can be selected as input. For algorithm, we divided the implementation into three stages. First is output is binary with shallow learning, and the second is multiclass output with shallow learning, and the last one is deep learning. The experiments is provided to compare the performance of the three stages.

1 Introduction

Intuitively, information usually correlate with other information in a certain range. One piece of information of a person may closely correlated with his/her other information. Besides, one person's information may also has some connection with the information of his/her friends. Based on this trait, it's plausible to use machine learning techniques to infer one's unknown information based on other related information. In fact, the utilization of machine learning technique on information analysis has gain increased popularity in recent years. For example, this paper (Melville, Prem, et al.,) gives a general summary of apply machine learning for social media analysis.

In this paper, we have dataset of a bunch of adults' information including marital status, education, salary level and etc. We implement a machine

learning system which can infer one feature of a instance based on it's other features.

Generally, there are three types of machine learning tasks, supervised learning, unsupervised learning and reinforcement learning. Since the training data we have includes input and output, the task in this paper is supervised learning. Based on the type of output, machine learning task can also be classified as classification, regression and clustering. In this paper, the output is either binomial or multinomial. So it is a classification task. A Neural Networks based algorithm can fulfil the task of supervised classification work. Besides, Neural networks are universal approximators, and they work well for capturing associations within a set of features, which is what we need in this task. Hence, we choose Multilayer Perceptron (MLP), which is one of feed-forward artificial neural network, to implement this task.

The original Multilayer Perceptron algorithm requires feature data to be discrete. However, sometimes some feature might be continuous in dataset, like the dataset we chose. In this case, we have to use an approach to transfer continuous data into discrete.

The rest of this paper is organized as follows. In section 2, we briefly describe the machine learning algorithm we choose. Section 3 presents the implementation details, which introducing dataset and 3 different levels of system we implement. In section 4 we provide experiment results and system performance under different conditions. At last, we conclude the paper in section 5.

2 Method

The algorithm implemented in this paper is Multi-layer Perceptron (MLP). MLP is an artificial neural network to solve classification problem. The bottom layer of the neural network is input layer which consists of several input nodes, and the number of input nodes equals the number of features. The top layer of the neural network is output layer. If the output only has two classifications, we can use only one output node, and it has binary value to represent two classifications. If the output has N , where $N > 2$, classifications, we need N nodes in output layer to represent the N classifications. The layers between input layer and output layer are hidden layers. The number of hidden layers and the nodes of each hidden layer are defined by users. The Fig. 1 represents the structure of MLP, as shown in the figure, you can see the input layer, hidden layer and output layer. It only shows one hidden layer, although you can add more than one hidden layer according to requirements.

2.1 Weighted Connection

Assume we have M layers totally, including input layer and output layer. Then, the net value of each node in layer $m + 1$ equals to the weighted summation of layer m .

$$net_j = \sum_{i \in l_m} w_{ij} net_i \quad (1)$$

where l_m stands for the nodes in layer m . As you can see in the Fig. 1, there are some arrows between two subsequent layers, which stand for the weighted connections. There is no connections between two nodes in the same layers. The net values of input layer is input feature, because they do not have previous layer.

2.2 Activation Function

For each node except the node in input layer, after getting the net value, an activation function is applied on the net value, and the result of the activation is the output of this node, which is used as the input of nodes in next layer. The activation function is a non-linear function, and in most cases, sigmoid function or hyperbolic tangent function is selected as

activation function.

$$o_j = g(net_j) = \frac{1}{1 + e^{-net_j}} \quad (2)$$

$$o_j = g(net_j) = \tanh(net_j) \quad (3)$$

2.3 Forward Propagation

Now, we have weighted connections between two layers and activation function on each node, so we can infer the next layer given the current layer. We have input features as the input layer, then we can do forward propagation layer by layer, and finally, we get the net value of nodes in output layer.

We use sigmoid function as activation function. For binary output, we only have one output node, and we have

$$\text{if } o_j = \frac{1}{1 + e^{-net_j}} < 0.5 \text{ then } Y = 0 \quad (4)$$

$$\text{if } o_j = \frac{1}{1 + e^{-net_j}} \geq 0.5 \text{ then } Y = 1 \quad (5)$$

For multiclass case which has M classifications, we encode the output with M nodes, and each node is still binary. For example, we have 4 classifications, we use four output nodes in output layer, and for class 1, the values of the 4 nodes are 1000, and for class 2, 4 nodes are 0100, and for class 3, 4 nodes are 0010, and for class 4, 4 nodes are 0001. We use the same method as that of binary class case to determine the binary value of each node in output layer when training the set. The predict methods in training step and testing accuracy is a little bit different for multiclass case. In the training stage, it might happen that more than one node are 1, because we use 0.5 as threshold for each sigmoid function, but we only have one node equals to one in real output, so when testing the accuracy, we only let the node which has the highest value equals to one.

2.4 Backward Propagation

The aim of training samples is to get appropriate weighted connections to predict accurate output. First, we use random weighted connections as initialization. Each step, an instance features are inputted and use current weighted connections to get a

predicted output, and use loss function

$$E = \frac{1}{2}(y - o_j)^2 \quad (6)$$

to measure the difference between predicted result and real output, where o_j means the output value of j^{th} node in output layer, and y is the real output. The partial derivatives of loss function on w_{ij} is

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial net_j} \frac{\partial net_j}{\partial w_{ij}} \quad (7)$$

According to equation 2, we have

$$\frac{\partial o_j}{\partial net_j} = g(net_j)(1 - g(net_j)) \quad (8)$$

According to equation 1, we have

$$\frac{\partial net_j}{\partial w_{ij}} = net_i \quad (9)$$

If node j is in output layer, according to equation 6, we have

$$\frac{\partial E}{\partial o_j} = o_j - y_j \quad (10)$$

Thus, we have

$$\frac{\partial E}{\partial w_{ij}} = (o_j - y_j)g(net_j)(1 - g(net_j))net_i$$

let $\frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial net_j} = \sigma_j$, we have

$$\frac{\partial E}{\partial w_{ij}} = \sigma_j net_i \quad (11)$$

If node j is not output layer, it is a little bit complicated.

$$\frac{\partial E}{\partial o_j} = \sum_{k \in K} \left(\frac{\partial E}{\partial net_k} \right) \left(\frac{\partial net_k}{\partial o_j} \right) = \sum_{k \in K} \left(\frac{\partial E}{\partial o_k} \right) \left(\frac{\partial o_k}{\partial net_k} \right) w_{jk} \quad (12)$$

Thus, we have

$$\frac{\partial E}{\partial w_{ij}} = \sigma_j net_i \quad (13)$$

where $\sigma_j = \sum_{k \in K} (\sigma_k w_{jk}) g(net_j)(1 - g(net_j))$

$$\Delta w_{ij} = \eta \frac{\partial E}{\partial w_{ij}} \quad (14)$$

where $\frac{\partial E}{\partial w_{ij}}$ is shown in equation 11 and 13, and η is learning rate.

2.5 Pre-training

Pre-training is used only when there are more than one hidden layers in the network. We tried two pre-training method, one is RBM, and another is contrastive divergence(CD-1). Both of them can be applied to binary input.

For RBM, it is introduced and implemented in homework4, so we just use the homework4 codes. We will not talk about it here. For CD-1, the weighted connections are undirected, which means you can use the weighted connections to get lower layer nodes value according to upper layer nodes. In general, CD-1 has three steps,

1. use nodes \mathbf{v}_n in lower layer n to get nodes \mathbf{v}_{n+1} in next layer $n + 1$ through current weighted matrix which consists of weighted connections.
2. use the nodes \mathbf{v}_{n+1} in next layer $n + 1$ to get nodes \mathbf{v}'_n in lower layer n through current weighted matrix.
3. use the nodes \mathbf{v}'_n in lower layer n to get nodes \mathbf{v}'_{n+1} in next layer $n + 1$ through current weighted matrix.

Now we have \mathbf{v}_n , \mathbf{v}_{n+1} , \mathbf{v}'_n and \mathbf{v}'_{n+1} , and we will use them to compute new \mathbf{w} , where \mathbf{w} is weighted matrix.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + a(\mathbf{v}_n \mathbf{v}_{n+1}^T - \mathbf{v}'_n \mathbf{v}_{n+1}'^T) \quad (15)$$

where a is learning rate.

The initialization of \mathbf{w} is randomized, and the pre-training result will be used as the initialization of training process.

3 Implementation

3.1 Dataset

We get the data set from UCI Machine Learning Repository. It contains many aspects of personal information. There are 32561 samples in training set and 16281 instances in testing set with unknown values, and if the unknown values are removed, training set has 30162 instances, and test set has 15060 instances. For binary case, we choose salary as output. If the *salary* < 50K, $Y = 0$, and if *salary* ≥ 50K, then $Y = 1$, following table shows the details about the personal features

Table 1: Data type of each feature, and the types after pre-processing.

Feature Name	Data Type	Converted Type
Age	continuous	continuous
Workclass	Category	Integer
Fnlwgt	continuous	continuous
Education	category	Integer
Education-num	continuous	continuous
Marital-status	category	Integer
Occupation	category	Integer
Relationship	category	Integer
Race	category	Integer
Sex	category	Integer
Capital-gain	continuous	continuous
Capital-loss	continuous	continuous
Hours-per-week	continuous	continuous
Native-country	category	Integer

For the multiclass case, the relationship is predicted according to other features and salary. For the category features, we assign different integer to each class from 1 to the total number of classification of that feature. After converting all the instance features to numerical, because the ranges are so different, we need to standardize these features

$$x_{ij} \leftarrow \frac{x_{ij} - \mu_i}{s_i} \quad (16)$$

where x_{ij} stands for the j^{th} feature of i^{th} instance, and μ_i stands for the mean of i^{th} feature and s_i means the standard deviation of i^{th} feature. All the standardization work can be done with *scale* function in R language.

3.2 Binary Output

We implement the MLP algorithm by Python. Training instances are stored in a text file, and the name of the file is an input parameter. We have only one hidden layer, so the user needs to input the number of the nodes in the hidden layer. We only have one node in the output layer, which is fixed, because it is a binary output.

First, we need read training instances into a list, which will be used in for loop. The weighted matrices between two layers are stored in a 3-dimension list, which is initialized as an random number between 0 and 1. For each instance, we will initialize the net value of each node to zero at first, and then

get a new net value after forward propagation, and all the net values are also stored in a 3-dimension list, and the list will be reset to zero when a new instance comes. In the backward propagation process, we need to store all the σ_i in the upper layer according to equation 13.

All the weighted matrices are updated when inputting an new instance, and after training all the instances, we can get two weighted matrices consists of weighted connections between input layer and hidden layer and those between hidden layer and output layer. The two weighted matrices are exported to a text file, which will be used to test the accuracy of training set and testing set.

For testing part, we need to input file which stores the weighted matrices given by the training process, then we use forward propagation to get predicted output, and compares with real output to get the accuracy.

3.3 Multiclass

Most part of multiclass is the same as binary output except the number of nodes in output layer and the difference between forward propagation and predicting process. The number of nodes in output layer equals to the number classifications, which needs to be provided by users. The forward propagation is the same as that in binary output case, but the predict process is implemented by another function to avoid more than one node in output layer equals to 1, which is impossible for real output. Thus only the node with highest value after sigmoid function equals to 1.

3.4 Deep Learning

The number of nodes in each hidden layer and output layer needs to be provided by users. We implement deep learning in three approaches. The first approach does not use pre-training and another two approaches use different pre-training method, one is RBM and another is CD-1.

For the first one which does not use pre-training, we only need to add more weighted matrices into the list which is used to store weight matrices, and other parts of the implementation is the same as multiclass.

For the two approaches which need pre-training, we use the first 500 samples in the training set as pre-

training set, and the rest of samples are still used as training set. Also, we need do some pre-process work on the input data, because RBM and CD-1 can be only applied on binary input. First, we use sigmoid function shown in equation 2 to make all the input features in the range $[0, 1]$, and use the converted value as probability in a Bernoulli trial to sample a value from $\{0, 1\}$, following is the Python code to do the pre-process work

```
bernoulli.rvs(1.0/(1+math.exp(-1*x)), size = 1)[0]
```

where x is the data before processing.

For implementation of RBM, we will not talk about it here, because it has been introduced in homework4. What we did is to export the result of RBM into a file, which can be used as the initialization of the weighted matrices.

For the implementation of CD-1, we use the 500 samples to pre-train the weighted matrix between input layer and the first hidden layer, then after getting the pre-trained weighted matrix, we use it to update the value of nodes in first hidden layer. We use the same method layer by layer until we get all the pre-trained weighted matrices between two neighboring layers. As for how to get one pre-trained weighted matrix, the three steps mentioned in 2.5 are implemented. All of the vectors \mathbf{v}_n , \mathbf{v}_{n+1} , \mathbf{v}'_n and \mathbf{v}'_{n+1} are binary, and the method to convert net value to binary is the same as the pre-process method, using sigmoid function and Bernoulli trial.

4 Experimental Results

We have conducted three experiments. Each experiment we test the system with different activation function and different numbers of hidden nodes.

4.1 Binomial Classification

The first system we implement can only infer binomial feature. The feature we choose as label is salary level, which only have two states, larger than 50k or otherwise. The number of feature used to training and predict is 14. We have trained and tested the data under 3 different number of hidden nodes: one that smaller than number of feature number, one that equal to the feature number and one that larger than feature number. The activation function used in this experiment is sigmoid. The training accuracy and test accuracy are shown in table 2.

Hidden Nodes Number	10	14	18
Training Accuracy	75.92	81.13	76.23
Test Accuracy	75.38	80.87	73.08

Table 2: Binary Accuracy Percentage

Hidden Nodes Number	10	14	18
Training Accuracy	70.49	80.81	83.91
Test Accuracy	70.23	79.77	83.81

Table 3: Multinomial(2 states) Accuracy Percentage, sigmoid

These results show that training accuracy and test accuracy will reach the highest value when the number of hidden nodes is the same as feature number.

4.2 Multinomial Classification

The improved system can infer feature with multinomial states. First we still choose the binomial feature we use in previous experiment as label. In addition, we build the system with different activation function, sigmoid and tangent, the results of which are shown in table 3 and table 4 respectively.

Then we choose — as feature which has 6 states. Results are shown in table 5 and table 6.

From these results we can tell that sigmoid activation function generally leads to better results than tangent activation function. Besides, when the number of output states increase, both the training and test accuracy will decrease.

4.3 Deep Network Binomial Classification

The final system use 2 hidden layer deep network learning algorithm. The target of this system is to test the effect of pre-training, so we only conduct the experiment with some fixed condition, that is, 14 hidden nodes on each hidden layer, binomial output and sigmoid activation function. The result is shown in table 7.

Hidden Nodes Number	10	14	18
Training Accuracy	75.92	81.13	76.23
Test Accuracy	75.38	80.87	73.08

Table 4: Multinomial(2 states) Accuracy Percentage, tangent

Hidden Nodes Number	10	14	18
Training Accuracy	75.80	75.98	73.12
Test Accuracy	74.23	74.84	72.30

Table 5: Multinomial(6 states) Accuracy Percentage, sigmoid

Hidden Nodes Number	10	14	18
Training Accuracy	45.53	51.35	51.72
Test Accuracy	40.27	49.63	52.06

Table 6: Multinomial(6 states) Accuracy Percentage, tangent

These results show that pre-training decrease the training accuracy and increase the test accuracy.

4.4 Experiments Summary

Conclude from the experimental results above, we can get following conclusions:

1. When the number of hidden nodes is less than the number of feature, the accuracies are definitely worse than when the number of hidden nodes is equal to the number of feature. But after that, the increase of hidden nodes number may not lead to higher accuracies. Usually, the use the number of data feature as number of nodes in input layer will give a decent performance. How to optimize the number of nodes in hidden layer is a far more complex problem and we will not discuss it in this paper.
2. In our experiments, sigmoid activation function generally produce better results than tangent function. This may due to that in our experiments, the activity in the network during training is close to positive 1 or negative 1. Then when using tanh, the gradient for the activation function may go to positive 1 or negative 1. Glorot and Bengio have discussed this problem in their paper (Glorot, X., Bengio, Y., 2010).
3. Unsupervised pre-training do help improve learn-

pre-training	yes	no
Training Accuracy	78.89	81.51
Test Accuracy	78.86	75.42

Table 7: Deep Multinomial(2 states) Accuracy Percentage, sigmoid

ing in deep networks. The reason why the training accuracy is lower in our third experiment may be that we don't have unsupervised data, so we have to use part of training data as unsupervised data which are therefore eliminate from training data. So in fact we less training data to do the learning. As a result, the training accuracy decreased.

4.5 Electronically-available resources

NAACL HLT provides this description in \LaTeX 2e (naaclhlt2010.tex) and PDF format (naaclhlt2010.pdf), along with the \LaTeX 2e style file used to format it (naaclhlt2010.sty) and an ACL bibliography style (naaclhlt2010.bst). These files are all available at <http://naaclhlt2010.isi.edu>. A Microsoft Word template file (naaclhlt2010.dot) is also available at the same URL. We strongly recommend the use of these style files, which have been appropriately tailored for the NAACL HLT 2010 proceedings.

4.6 Format of Electronic Manuscript

For the production of the electronic manuscript you must use Adobe's Portable Document Format (PDF). This format can be generated from postscript files: on Unix systems, you can use ps2pdf for this purpose; under Microsoft Windows, you can use Adobe's Distiller, or if you have cygwin installed, you can use dvipdf or ps2pdf. Note that some word processing programs generate PDF which may not include all the necessary fonts (esp. tree diagrams, symbols). When you print or create the PDF file, there is usually an option in your printer setup to include none, all or just non-standard fonts. Please make sure that you select the option of including ALL the fonts. *Before sending it, test your PDF by printing it from a computer different from the one where it was created.* Moreover, some word processor may generate very large postscript/PDF files, where each page is rendered as an image. Such images may reproduce poorly. In this case, try alternative ways to obtain the postscript and/or PDF. One way on some systems is to install a driver for a postscript printer, send your document to the printer specifying "Output to a file", then convert the file to PDF.

For reasons of uniformity, Adobe's **Times Roman** font should be used. In L^AT_EX2e this is accomplished by putting

```
\usepackage{times}  
\usepackage{latexsym}
```

in the preamble.

Additionally, it is of utmost importance to specify the **US-Letter format** (8.5in × 11in) when formatting the paper. When working with dvips, for instance, one should specify `-t letter`.

Print-outs of the PDF file on US-Letter paper should be identical to the hardcopy version. If you cannot meet the above requirements about the production of your electronic submission, please contact the publication chairs above as soon as possible.

4.7 Layout

Format manuscripts two columns to a page, in the manner these instructions are formatted. The exact dimensions for a page on US-letter paper are:

- Left and right margins: 1in
- Top margin: 1in
- Bottom margin: 1in
- Column width: 3.15in
- Column height: 9in
- Gap between columns: 0.2in

Papers should not be submitted on any other paper size. Exceptionally, authors for whom it is *impossible* to format on US-Letter paper, may format for A4 paper. In this case, they should keep the *top* and *left* margins as given above, use the same column width, height and gap, and modify the bottom and right margins as necessary. Note that the text will no longer be centered.

4.8 The First Page

Center the title, author's name(s) and affiliation(s) across both columns. Do not use footnotes for affiliations. Do not include the paper ID number assigned during the submission process. Use the two-column format only when you begin the abstract.

Title: Place the title centered at the top of the first page, in a 15 point bold font. (For a complete guide to font sizes and styles, see Table 8.) Long title should be typed on two lines without a blank line intervening. Approximately, put the title at 1in from the top of the page, followed by a blank line, then the author's names(s), and the affiliation on the following line. Do not use only initials for given names (middle initials are allowed). Do not format surnames in all capitals (e.g., "Leacock," not "LEA-COCK"). The affiliation should contain the author's complete address, and if possible an electronic mail address. Leave about 0.75in between the affiliation and the body of the first page.

Abstract: Type the abstract at the beginning of the first column. The width of the abstract text should be smaller than the width of the columns for the text in the body of the paper by about 0.25in on each side. Center the word **Abstract** in a 12 point bold font above the body of the abstract. The abstract should be a concise summary of the general thesis and conclusions of the paper. It should be no longer than 200 words. The abstract text should be in 10 point font.

Text: Begin typing the main body of the text immediately after the abstract, observing the two-column format as shown in the present document. Do not include page numbers.

Indent when starting a new paragraph. For reasons of uniformity, use Adobe's **Times Roman** fonts, with 11 points for text and subsection headings, 12 points for section headings and 15 points for the title. If Times Roman is unavailable, use **Computer Modern Roman** (L^AT_EX2e's default; see section 4.6 above). Note that the latter is about 10% less dense than Adobe's Times Roman font.

4.9 Sections

Headings: Type and label section and subsection headings in the style shown on the present document. Use numbered sections (Arabic numerals) in order to facilitate cross references. Number subsections with the section number and the subsection number separated by a dot, in Arabic numerals.

Citations: Citations within the text appear in parentheses as (Gusfield, 1997) or, if the author's name appears in the text itself, as Gusfield (1997). Append lowercase letters to the year in cases of am-

biguities. Treat double authors as in (Aho and Ullman, 1972), but write as in (Chandra et al., 1981) when more than two authors are involved. Collapse multiple citations as in (Gusfield, 1997; Aho and Ullman, 1972).

References: Gather the full set of references together under the heading **References**; place the section before any Appendices, unless they contain references. Arrange the references alphabetically by first author, rather than by order of occurrence in the text. Provide as complete a citation as possible, using a consistent format, such as the one for *Computational Linguistics* or the one in the *Publication Manual of the American Psychological Association* (American Psychological Association, 1983). Use of full names for authors rather than initials is preferred. A list of abbreviations for common computer science journals can be found in the *ACM Computing Reviews* (Association for Computing Machinery, 1983).

The L^AT_EX and BibT_EX style files provided roughly fit the American Psychological Association format, allowing regular citations, short citations and multiple citations as described above.

Appendices: Appendices, if any, directly follow the text and the references (but see above). Letter them in sequence and provide an informative title: **Appendix A. Title of Appendix.**

Acknowledgment sections should go as a last (unnumbered) section immediately before the references.

4.10 Footnotes

Footnotes: Put footnotes at the bottom of the page. They may be numbered or referred to by asterisks or other symbols.¹ Footnotes should be separated from the text by a line.² Footnotes should be in 9 point font.

4.11 Graphics

Illustrations: Place figures, tables, and photographs in the paper near where they are first discussed, rather than at the end, if possible. Wide illustrations may run across both columns and should be placed at the top of a page. Color illustrations are

¹This is how a footnote should appear.

²Note the line separating the footnotes from the text.

Type of Text	Font Size	Style
paper title	15 pt	bold
author names	12 pt	bold
author affiliation	12 pt	
the word “Abstract”	12 pt	bold
section titles	12 pt	bold
document text	11 pt	
abstract text	10 pt	
captions	10 pt	
bibliography	10 pt	
footnotes	9 pt	

Table 8: Font guide.

discouraged, unless you have verified that they will be understandable when printed in black ink.

Captions: Provide a caption for every illustration; number each one sequentially in the form: “Figure 1. Caption of the Figure.” “Table 1. Caption of the Table.” Type the captions of the figures and tables below the body, using 10 point text.

5 Length of Submission

The NAACL HLT 2010 main conference accepts submissions of long papers and short papers. The maximum length of a long paper manuscript is eight (8) pages of content and one (1) additional page of references *only* (appendices count against the eight pages, not the additional one page). The maximum length of a short paper manuscript is four (4) pages including references. For both long and short papers, all illustrations, references, and appendices must be accommodated within these page limits, observing the formatting instructions given in the present document. Papers that do not conform to the specified length and formatting requirements are subject to be rejected without review.

Acknowledgments

Do not number the acknowledgment section.

References

Melville, Prem, Vikas Sindhwani, Richard D. Lawrence, Estepan Meliksetian, Yan Liu, Pei-Yun Hsueh, and Claudia Perlich. *Machine Learning for Social Media Analytics*

- Glorot, Xavier, and Yoshua Bengio. 2010. *Understanding the difficulty of training deep feedforward neural networks*. In *International Conference on Artificial Intelligence and Statistics*, pp. 249-256.
- Alfred V. Aho and Jeffrey D. Ullman. 1972. *The Theory of Parsing, Translation and Compiling*, volume 1. Prentice-Hall, Englewood Cliffs, NJ.
- American Psychological Association. 1983. *Publications Manual*. American Psychological Association, Washington, DC.
- Association for Computing Machinery. 1983. *Computing Reviews*, 24(11):503–512.
- Ashok K. Chandra, Dexter C. Kozen, and Larry J. Stockmeyer. 1981. Alternation. *Journal of the Association for Computing Machinery*, 28(1):114–133.
- Dan Gusfield. 1997. *Algorithms on Strings, Trees and Sequences*. Cambridge University Press, Cambridge, UK.




Figure 1: Structure of MLP with one hidden layer